

# CS340 A2


## User space file systems

# File Systems

File systems are simply different ways of **organizing** and **storing files** on a **storage device**.

It needs to satisfy these general requirements.

- We need some way of storing information
  - independently from a running program
  - so it can be used at a later time
  - permanently (or an approximation to it)
  - non-volatile storage
  - so it can be shared with other programs or users
- An infinite variety of data is to be stored
  - The more information the OS knows about the data the more it can facilitate use of the data.
  - e.g. Executable files
- Naming the data
  - The data needs to be stored and retrieved easily. We need a way to name the data.
  - We must then be able to locate the data using its name.



# File Operations

On most systems these commands need security authorisation to perform and they work on the file as a whole.

## Create

- Need to specify information about the file:
  - the name
  - the file type (or some representation of the program associated with this file)
- Do we need to specify the size of the file? (certainly helps with keeping storage contiguous but is usually regarded as an unnecessary restriction)
- Creation needs to do something to the associated device - at least write to some structure (sometimes a directory).
- Some systems allow transitory files to not be recorded permanently in secondary storage.

# File Operations

## Delete

- Remove the file (or the directory entry). Release the space used by the file, so it could be reused by other files.

## Move

- Moving a file can be performed in different ways depending on the before and after locations.
- If both locations are on the same device the data doesn't have to be copied and then the original deleted. Instead change information about the file.

## Open

- The OS evaluates the name, check the access permission.

# File Operations

## Read/Write

These operations work on the file's contents.


- We need to know where the information is on the device.
- Must specify what data to read/write, how much, and where to put it.
- Write requires the allocation of extra space.

# File system in USEr space - FUSE

- Linux has a library which allows file systems to be written and used without root privileges - libfuse.
- The library is used in conjunction with a kernel module which provides the privileged operations.
- The user's file system gets mounted on to an existing directory.


# Mounting a file system

- A file system (in Unix context) can be thought of as a device or partition which can be connected into the standard hierarchical file system starting from the root `/`.

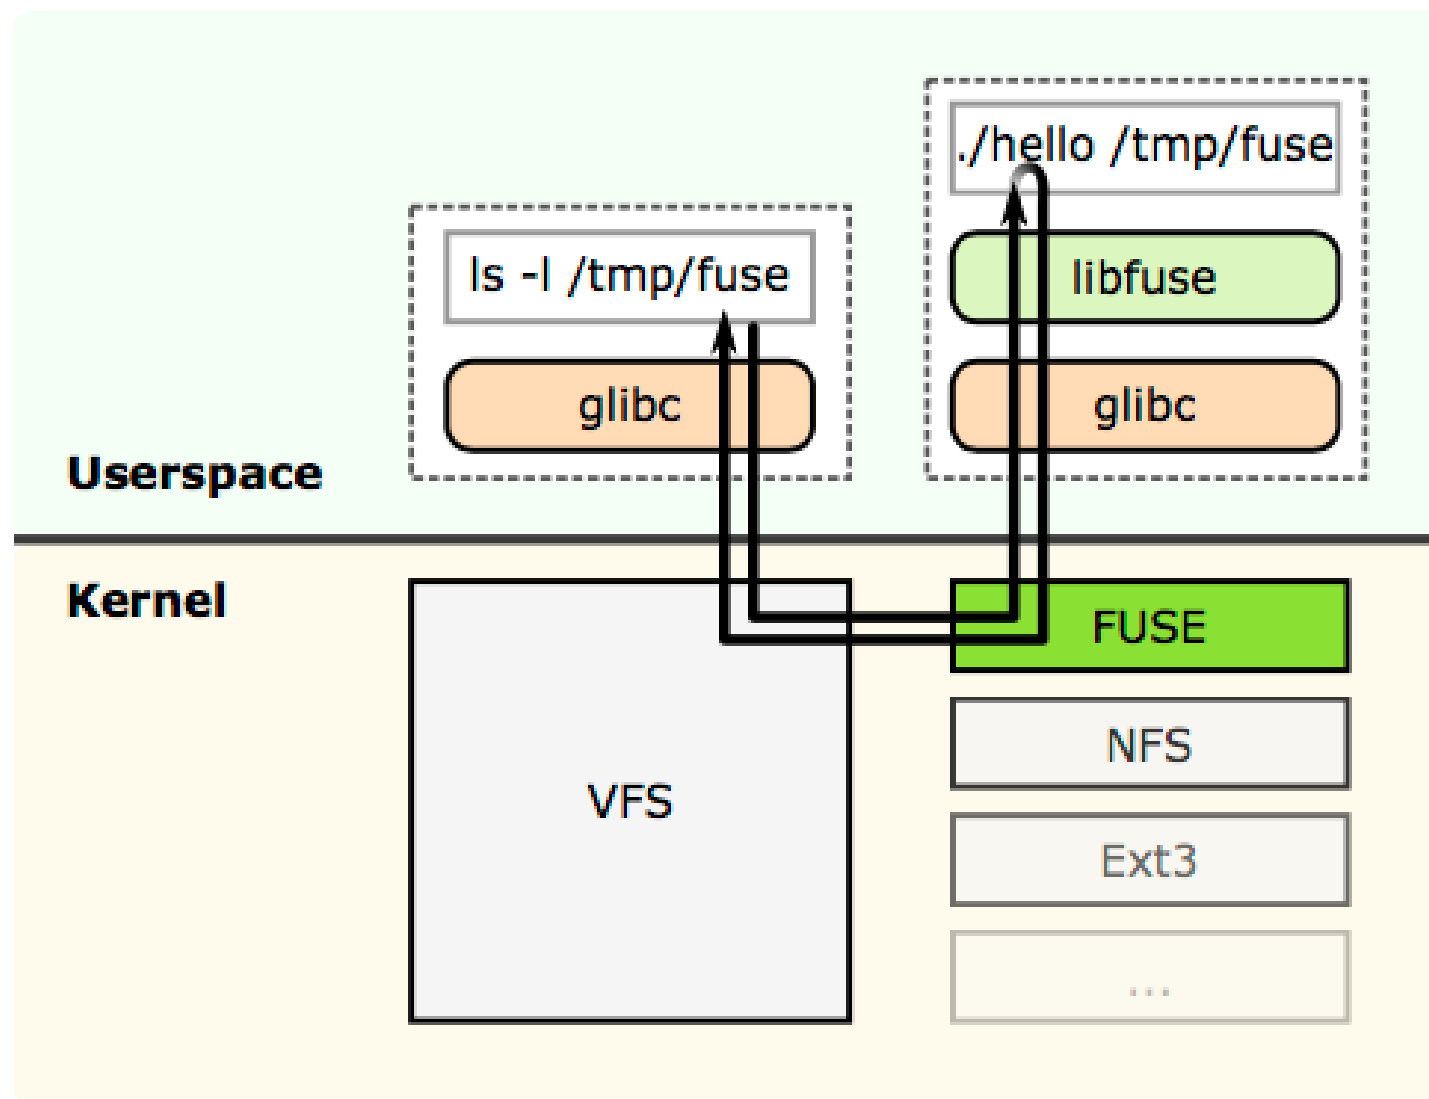


# Mounting a file system

- We need a **mount point**: a directory where we plant the file system
  - e.g., /exp/home
- Any files which are already in the mount point directory then disappear
- They get replaced by the files/directories in the new file system







# How it works

from Wikipedia

`./hello` is mounted on `/tmp/fuse`

All the operations on `/tmp/fuse` and its files can be handled by user level code.

# How it works

- Need a handler program which gets linked to the FUSE library.
- The handler program specifies how the files respond to read/write/stat requests.
- At the time of mounting the new file system (handler gets registered with the kernel).
- Mainly useful for virtual file systems.

# Setup

- Ubuntu in the labs or on your own machine.
  - Virtual machines, macOS or Windows Subsystem for Linux version 2
  - Ubuntu 18.04 image has been made available at FlexIT ([flexit.auckland.ac.nz](https://flexit.auckland.ac.nz))
    - Sign in to SSO which will take you to the sign-in page for FlexIT
    - Look for the Ubuntu 18.04 desktop icon
- The Markers will use Ubuntu in the labs.

# Setup

- Download fuse.py, memory.py and passthrough.py from the A2 files section on Canvas.
  - fuse.py is from
    - <https://github.com/fusepy/fusepy/blob/master/fuse.py>
  - memory.py is from
    - <https://github.com/fusepy/fusepy/blob/master/examples/memory.py>
  - passthrough.py is from
    - <https://github.com/skorokithakis/python-fuse-sample>
- You can also get examples from the GitHub site at
  - <https://github.com/fusepy/fusepy/tree/master/examples>

# OMG Python

- Much (much) simpler to do this assignment in Python.
- If you don't know how to do something in Python.
  - Ask one of the tutors/lecturers for help
  - Forum: feel free to ask on Piazza or Canvas
  - Python documentation: <https://docs.python.org/3.6/>
  - StackOverflow, Google
- The standard python 2 or 3.6 on Ubuntu works.

# passthrough.py

- Neatly encapsulates the methods you may need to override (not all of them)
- Two types of methods
  1. filesystem methods which deal with directories and files
  2. file methods which deal with the contents of files

# Logging

- There is some python magic which logs information to the screen as the file system works.

```
class LoggingMixin:
    log = logging.getLogger('A2')

    def __call__(self, op, path, *args):
        self.log.debug('-> path:%s %s%s', path, op, repr(args))
        ret = '[Unhandled Exception]'
        try:
            ret = getattr(self, op)(path, *args)
            return ret
        except OSError as e:
            ret = str(e)
            raise
```

# Part1

- Setup
  - Put `fuse.py`, `passthrough.py` and `a2fuse1.py` in the same directory
  - Create a directory called "source"
  - Download files "one", "two", and "three" from Canvas into the "source" directory.
  - Create a directory called "mount"
- Run `python a2fuse1.py source mount`
  - Observe logs and highlight the main functionality, for example

```
DEBUG:fuse.log-mixin:-> getattr /newfile (None,)
DEBUG:fuse.log-mixin:<- getattr {'st_atime': 1599037397.463398, 'st_ctime':
1599037379.543504, 'st_gid': 20, 'st_mode': 33188, 'st_mtime':
1599037379.543504, 'st_nlink': 1, 'st_size': 12, 'st_uid': 501}
```

Gets attributes for "newfile".

- To stop
  - `fusermount -u mount` (Ubuntu) or `umount mount` (macOS)



# Part2

- In the Memory class, explain what each method does
  - `__init__`, `getattr`, `readdir`, `open`, `create`, `unlink`, `write`, `read`

```
def __init__(self):  
    self.files = {}  
    self.data = defaultdict(bytes)  
    self.fd = 0  
    now = time()  
    self.files['/'] = dict(st_mode=(S_IFDIR | 0o755), st_ctime=now,  
                          st_mtime=now, st_atime=now, st_nlink=2)
```

- Example: `__init__`
  - Creates an empty dictionary `self.files` for the files, using the path names as the keys. Each value in the dictionary will be another dictionary.
  - `self.data` is a dictionary for the files' data. The path names are the keys, the values are the data of that file.
  - Sets the starting value for the file descriptors, these are going to be used as unique file identifiers.
  - Grabs the current time and sets the file attributes for the root of this file system. It is a directory, with creation, modified and accessed times set to now. It has two links.

# Part3 – Requirement 1

- Setup
  - Create two directories, `source1` and `source2`
  - Put file `one` in `source1`
  - Put file `two` in `source2`
  - Make sure that `mount` is initially empty
- Enable your FUSE to mount two source directories into a mount point
  - Override the “`_full_path`” method to add logic to return multiple source directories
  - Override the “`readdir`” function to add logic s.t.
    - “`ls`” command in your `mount` can list the content of both `source1` and `source2`
  - Override the “`main`” method because we need an extra parameter
- Run “`python a2fuse2.py source1 source2 mount`”

# Part3 – Requirement 2

- Create your own user space file system
- The file system has two classes of files in the `mount` directory. One consists of real files from the `source1` and `source2` directories and the other of files which only exist in memory.
  - Override the File methods to create a totally cached file system
  - When a file is opened read all of the contents into memory
  - Any reads are extracted directly from memory with no disk access
  - Any writes go to memory
  - When the file is closed the changes get written to disk

# Submission

- Use the Canvas submission system to submit your assignment
- Zip together [A2.txt](#) and [a2fuse2.py](#)
- Extra marks
  - 1 mark for including your name and login in both files
  - 1 mark for any files created by the file system having the correct user and group ids